

Probabilistic Logic Programming

Xin Zhang

Peking University

Recap of Last Lecture

- $(\mathbf{S}, \mathbf{B})$ is a measurable space
 - $\mathbf{S}$ is a set
 - $\mathbf{B}$ is a σ -algebra on $\mathbf{S}$, which is a collection of subsets of $\mathbf{S}$
 - It contains $\emptyset$
 - Closed under complementation in $\mathbf{S}$
 - Closed under countable union
 - The elements of $\mathbf{B}$ are called measurable sets
- A signed (finite) measure on $(\mathbf{S}, \mathbf{B})$ is a countably additive map $\mu: \mathbf{B} \rightarrow \mathbf{R}$ such that $\mu(\emptyset) = 0$

Recap of Last Lecture

- Given $f: (\mathbf{S}, \mathbf{B}_\mathbf{S}) \rightarrow (\mathbf{T}, \mathbf{B}_\mathbf{T})$ measurable, a measure μ on $\mathbf{B}_\mathbf{S}$, the **pushforward measure** $\mu(f^{-1}(B))$ on $\mathbf{B}_\mathbf{T}$ is defined as

$$f_*(\mu)(B) = \mu(f^{-1}(B)), \quad B \in \mathcal{B}_T.$$

- Given $(\mathbf{S}, \mathbf{B}_\mathbf{S})$ and $(\mathbf{T}, \mathbf{B}_\mathbf{T})$, $P: \mathbf{S} \times \mathbf{B}_\mathbf{T} \rightarrow \mathbf{R}$ is called a Markov kernel if
 - For fixed $A \in \mathbf{B}_\mathbf{T}$, the map $\lambda s. P(s, A) \rightarrow \mathbf{R}$ is a measurable function on $(\mathbf{S}, \mathbf{B}_\mathbf{S})$
 - For fixed $s \in \mathbf{S}$, the map $\lambda A. P(s, A) \rightarrow \mathbf{R}$ is a probability measure on $(\mathbf{T}, \mathbf{B}_\mathbf{T})$

Recap of Last Lecture

- Program semantics
 - Operational
 - Denotational
 - Measure theory

This Lecture

- Probabilistic logic programming
 - Motivation
 - Syntax
 - Semantics
 - Inference

Classical AI: Logic

- Rich logic systems provide significant expressiveness power
 - Concise and learnable models
- Example: first-order logic. Rules of chess occupy
 - 10^0 pages of first-order logic
 - 10^5 pages in propositional logic
 - 10^{38} pages in finite automata

Quick Recap on First-Order Logic

- Compared to propositional logic, introduces predicates and quantifications for expressiveness

$$\forall h1, h2, h3. sibling(h1, h2) \wedge sibling(h2, h3) \rightarrow sibling(h1, h3)$$

- Undecidable

Modern AI: Probability Theory for Uncertainty

- Bayesian network
- Fixed variables in fixed ranges
 - Similar to propositional logic and Boolean logic

Probabilistic Logic Programming: Unifying Logic and Probability


- Logic: the ability to describe complex domains concisely in terms of objects and relations
- Probability: the ability to handle uncertainty
- Logic + probability = Probabilistic Logic Programming

Example Probabilistic Logic Languages

- Markov Logic Network. University of Washington
- Probabilistic Soft Logic. University of Maryland
- **Problog**. KU Leuven <https://dtai.cs.kuleuven.be/problog/index.html>
- BLOG. UC Berkeley
-

Background: Logic Programming

- Declarative: specifies what rather than how
- Leverages powerful inference engine



Background: Prolog and Datalog

- Prolog: once popular in AI, still being used in pattern matching (NLP)
 - Turing-complete
- Datalog: a subset of Prolog
 - Can only express polynomial algorithms
 - Originates from the Database community (SQL with recursions)
 - Logic part of Prolog

Background: Datalog

Input Relation:

Edge(e1, e2)

Output Relation:

Path(e1, e2)

Rules:

Path(e1, e2) :- edge(e1, e2)

$\forall \text{edge}(e_1, e_2) \Rightarrow \text{path}(e_1, e_2)$

Path(e1, e3) :- path(e1, e2), edge(e2, e3)

$\forall \text{path}(e_1, e_2) \wedge \text{edge}(e_2, e_3) \Rightarrow \text{path}(e_1, e_3)$

Background: Datalog

Edge(1, 2) Edge(2, 3)

Path(1, 2) :- Edge(1, 2)

Path(2, 3) :- Edge(2, 3)

Path(1, 3) :- Path(1, 2), Edge(2, 3)



Adding Probabilities to Datalog

- If A is a friend of B , and B is a friend of C , then A is likely a friend of C .

Can you write a program for the above sentence?

Adding Probabilities to Datalog

- Suppose edges exist with probabilities (by observation), compute path reachability.

Can you write a program for the above sentence?

Add probabilities to rules or facts?

What is the semantics?

$\text{Path}(E1, E2) \text{ :- edge}(E1, E2)$
0.5: $\text{Path}(E1, E3) \text{ :- path}(E1, E2), \text{edge}(E2, E3)$

Given a set of derived tuples/facts, assign a probability to them.

Problog: Introduction

- A language developed by the group led by Luc De Raedt at KU Leuven
- Extends Prolog with probabilities
 - Actually closer to Datalog

Problog: Syntax

- **Value:** numbers, mixed numbers and letters starting with a letter in lower cases
- **Variable:** starting with a capital letter

Problog: Syntax

Definition	Example
fact	<code>a.</code>
probabilistic fact	<code>0.5::a.</code>
clause	<code>a :- x.</code>
probabilistic clause	<code>0.5::a :- x.</code>
annotated disjunction	<code>0.5::a; 0.5::b.</code>
annotated disjunction	<code>0.5::a; 0.5::b :- x.</code>

From the documentation of Problog

Problog: Syntax

Queries:

```
0.5::heads(C).  
two_heads :- heads(c1), heads(c2).  
query(two_heads).
```

```
0.5::heads(C) :- between(1, 4, C).  
query(heads(C)).
```

Evidence:

```
0.5::heads(C).  
two_heads :- heads(c1), heads(c2).  
evidence(\+ two_heads).  
query(heads(c1)).
```

From the documentation of Problog

Example Program I

0.5 :: stayUp.

0.7 :: drinkCoffee :- stayUp.

0.5 :: drinkCoffee :- \+ stayUp.

0.9 :: fallSleep :- \+ drinkCoffee, stayUp.

0.3 :: fallSleep :- drinkCoffee, stayUp.

0.1 :: fallSleep :- \+ stayUp.

evidence(fallSleep).

query(stayUp).

What does the following program compute?

0.5 :: stayUp.

0.7 :: drinkCoffee :- stayUp.

0.5 :: drinkCoffee :- \+ stayUp.

0.9 :: fallSleep :- \+ drinkCoffee, stayUp.

0.3 :: fallSleep :- drinkCoffee, stayUp.

0.1 :: fallSleep :- \+ stayUp.

query(stayUp).

evidence(fallSleep).

What does the following program compute?

0.5::heads1.

0.5::heads2.

0.6::heads1 :- heads2.

query(heads1).

query(heads2).

What does the following program compute?

0.5::heads1.

0.5::heads2.

\+ heads1 :- heads2.

query(heads1).

query(heads2).

Example Program 2

0.9 :: edge(0,1).

0.8 :: edge(1,2).

0.7 :: edge(2,3).

0.8 :: edge(2,4).

1 :: path(A,B) :- edge(A,B).

0.8 :: path(A,C) :- path(A,B), edge(B,C).

evidence(\+ path(0,3)).

query(path(0,4)).

Semantics of Problog

- What is the semantics of the following program?

0.5 :: stayUp.

0.7 :: drinkCoffee :- stayUp.

0.3 :: fallSleep :- drinkCoffee, stayUp.

query(fallSleep).

Semantics of Problog


- For simplicity, we assume all probabilities are attached to facts
- First idea: we can convert the program into a Bayesian network, but how?

Semantics of Problog

- Converting into a Bayesian network is viable, but there are small catches
- We give another semantics that defines a distribution of Datalog programs

Semantics of Problog

- From a Problog program, we can sample a Datalog program by sampling the facts



Semantics of Problog

- What about queries?

0.5 :: stayUp.

0.7 :: r1.

0.3 :: r2.

drinkCoffee :- stayUp, r1.

fallSleep :- drinkCoffee, stayUp, r2.

query(fallSleep)

A query calculates a marginal probability of a fact. Informally,

$$p(f) = \frac{\sum p(\text{any program that derives } f)}{\sum p(\text{any program})}$$

Semantics of Problog

- What about evidence?

```
0.5 :: stayUp.  
0.7 :: r1.  
0.3 :: r2.  
drinkCoffee :- stayUp, r1.  
fallSleep :- drinkCoffee, stayUp, r2.  
  
evidence(\+ fallSleep)  
query(stayUp)
```

Evidence filters out certain programs. Informally,

$$p(f) = \frac{\sum p(\text{any program that derives } f | \text{evidence})}{\sum p(\text{any program} | \text{evidence})}$$

Semantics of Problog

- What about relations and quantified variables?

0.9 :: edge(0,1).

0.8 :: edge(1,2).

0.7 :: edge(2,3).

0.8 :: edge(2,4).

path(A,B) :- edge(A,B).

0.8 :: path(A,C) :- path(A,B), edge(B,C).

evidence(\+ path(0,3)).

query(path(0,4)).

Semantics of Problog

- Move probabilities to facts

0.9 :: edge(0,1).

0.8 :: edge(1,2).

0.7 :: edge(2,3).

0.8 :: edge(2,4).

0.8 :: r(A,B,C).

path(A,B) :- edge(A,B).

path(A,C) :- path(A,B), edge(B,C), r(A,B,C).

evidence(\+ path(0,3)).

query(path(0,4)).

Semantics of Problog

- Ground

Constants: 0, 1, 2, 3, 4

$\text{path}(A,C) \text{ :- path}(A,B), \text{edge}(B,C), \text{r}(A,B,C).$

Generates

$\text{path}(0,0) \text{ :- path}(0,0), \text{edge}(0,0), \text{r}(0,0,0).$

$A=0, B=0, C=0$

$\text{path}(0,1) \text{ :- path}(0,0), \text{edge}(0,1), \text{r}(0,0,1).$

$A=0, B=0, C=1$

$\text{path}(0,1) \text{ :- path}(0,0), \text{edge}(0,1), \text{r}(0,0,1).$

$A=0, B=0, C=1$

...

Semantics of Problog

- After grounding, each ground term can be seen as a Boolean variable, then the whole program can be solved using the semantics of the Boolean case

$\text{path}(0,0) \rightarrow t1, \text{edge}(0,0) \rightarrow t2, r(0,0,0) \rightarrow t3$

$\text{path}(0,0) \text{ :- } \text{path}(0,0), \text{edge}(0,0), r(0,0,0).$



$t1 \text{ :- } t1, t2, t3$

Semantics of Problog

- First, ground the program into a Boolean program
- The Boolean program describes a distribution of Datalog program, which in turn defines a distribution of outputs

Questions

- Can you use Problog to express uniform distributions?
- What about loops?

Logic Part in Problog is more than Datalog

```
:- use_module(library(aggregate)).
```

```
pull(0).
```

```
count(1).
```

```
pull(N+1) :- pull(N), N < 10.
```

```
0.1 :: pull_SSR(N) :- pull(N).
```

```
num_SSRs(sum<X>) :- pull_SSR(N),count(X).
```

```
query(num_SSRs(X)).
```

But It is also Not Prolog

- The following program terminates in Problog but not in Prolog

```
child(anne,bridget).
```

```
child(bridget,caroline).
```

```
child(caroline,donna).
```

```
child(donna,emily).
```

```
descend(X,Y) :- descend(Z,Y), child(X,Z).
```

```
descend(X,Y) :- child(X,Y).
```

```
query(descend(anne,emily))
```


Inference

- As described before, inference can be done in two steps:
 - **Grounding.** Convert the program into a probabilistic program with only Boolean variables (no quantifiers)
 - **Solving.** Solve with the Boolean program produced above.

Optimization on Grounding

- Grounding replaces all variables with their values
 - Number of grounded rules is proportional to cartesian product of the domain sizes
- How to optimize?
 - A simple idea: only ground the part that is relevant to the queries and evidence.
 - Backtrack over the rules starting from the queries and evidence (SLD resolution).
 - A further optimization: stop tracking if a rule body doesn't hold according to the evidence

Optimization on Grounding

- If the logic part is Datalog without negation, we can use a Datalog solver to compute the grounding
- Datalog without negation is monotonic: the more rules or input facts, the more output facts
- If negation is on the input, it is still fine

Negation in Prolog

- Unfortunately, Prolog allows the following program:

one(1).

odd(X) :- one(X).

even(X) :- \+ odd(X).

And

0.5::a. 0.9 :: e:-a.

0.5::b. 0.9 :: e:-b.

0.1 :: \+e:-a,b

If such negations are not present, we can use a Datalog solver to ground, which is highly efficient.

Solving

- Once we have a grounded program, we can leverage existing techniques
- Idea 1: convert the program into a Bayesian network
- Idea 2: convert the program into a Boolean formula with weights (MaxSAT)

Solving: Converting into a Bayesian Net

0.8 :: a.

0.7 :: b.

0.5 :: c:- a.

0.5 :: c:- b.

query(c).

Solving: Converting into a Bayesian Net

- We move all probabilities to input facts
- We add a root node whose prior distribution is $P(r = 1) = 1$. Then we add a rule $p :: f:-r$ for each input fact $p::f$
- For each fact f , suppose it is derived using $r_1, \dots, r_n$, we add arcs from all facts in the rule bodies to f .
- We set conditional probabilities:
$$p(f \mid \vee body(r_i) == True) = 1$$
$$p(f \mid \vee body(r_i) == False) = 0$$

Only works for program without cycles

Solving: Converting into a MaxSAT

- Finding the most likely solution becomes solving the MaxSAT
- Computing marginal probabilities becomes weighted model counting

Brief Introduction on MaxSAT

MaxSAT:

	$a \wedge$	(C1)		Subject to	C1
	$\neg a \vee b \wedge$	(C2)			C2
4	$\neg b \vee c \wedge$	(C3)	=		
2	$\neg c \vee d \wedge$	(C4)		Maximize	4
7	$\neg d$	(C5)			× C3 + 2 × C4 + 7 × C5

Solution: $a = \text{true}, b = \text{true}, c = \text{true}, d = \text{false}$
 (Objective = 11)

Brief Introduction on MaxSAT

- Popular MaxSAT solving techniques: converting the problem into a series of SAT problem
- Brief idea: can any solution satisfy k clauses?
 - Linear search
 - Binary search
 - (UNSAT) core guided

Core-Guided MaxSAT Solving

- UNSAT core: a set of clauses which are not satisfiable
 - Minimum UNSAT core: removing any clause will make it satisfiable
 - Modern SAT solvers come with the ability to return UNSAT cores
- [Fu & Malik]: Each time allow one and only one clause to be relaxed

Example using MaxSAT for Inference

0.6 :: rain.

0.5 :: sprinkle.

0.9 :: grass_wet :- rain, sprinkle.



ln0.6 rain

ln0.4 !rain

ln0.5 sprinkle

ln0.5 !sprinkle

ln0.9 r

ln0.1 !r

grass_wet or !rain or !sprinkle or !r

!grass_wet or rain

!grass_wet or sprinle

!grass_wet or r

grass_wet :- rain, sprinkle is translated into
 $grass_{wet} \leftrightarrow rain \wedge sprinkle \wedge r$

Example using MaxSAT for Inference

- When translating rules, we have to consider the least fixed point semantics of Datalog
- Suppose the rules are acyclic, for a given fact f , we have to consider all grounded rules that derive f

$$f \leftrightarrow \bigvee body(r_i)$$

Example using MaxSAT for Inference

- When rules are cyclic, problems become complicated:

0.5::a. b:-a. b:-c. c:-b

- For reference:
 - Janhunen, T. 2004. Representing normal programs with clauses. In In Proc. of the 16th European Conference on Artificial Intelligence. IOS Press, 358–362.
 - Mantadelis, T. and Janssens, G. 2010. Dedicated tabling for a probabilistic setting. In Tech. Comm. of 26th International Conf. on Logic Programming. 124–133.

Brief Introduction on Weighted Model Counting

- Model counting: compute the number of assignments to a SAT expression

a or b 3 assignments

- Weighted model counting
 - Each variable has a weight for each assignment: $w(v)$
 - The model weight is the the product of variable weights
 - Now the count is a weighted sum

Example using WMC for Inference

0.6 rain

0.4 !rain

0.5 sprinkle

0.5 !sprinkle

$$w(\text{rain} = \text{true}) = 0.6$$

$$w(\text{rain} = \text{false}) = 0.4$$

$$w(\text{sprinkle} = \text{true}) = 0.5$$

$$w(\text{sprinkle} = \text{false}) = 0.5$$

grass_wet or !rain or !sprinkle

$$P(\text{grass_wet} = \text{true}) = \text{WMC}(M \wedge \text{grass_wet} = \text{true})$$

!grass_wet or rain

!grass_wet or sprinle

What if we want to evaluate

$$P(\text{rain} \mid \text{grass_wet} = \text{true})?$$

Using WMC for Marginal Inference

- Let the constructed weighted formula be M , queries be Q , evidence be E , then

$$P(Q) = \frac{WMC(M \wedge Q \wedge E)}{WMC(M \wedge E)}$$

- For more, refer to

Sang, T., Beame, P. and Kautz, H., 2005. Solving Bayesian networks by weighted model counting. In Proceedings of the Twentieth National Conference on Artificial Intelligence (AAAI-05) (Vol. 1, pp. 475-482). AAAI Press.

Further Reading on Problog

- <https://dtai.cs.kuleuven.be/problog/index.html>

Next Lecture

- Causality